

Measuring Learning in Informative Processes

Robert C. Luskin, Ariel Helfer, and Gaurav Sood

Working draft

Abstract

We often want to know how much people learn from informative processes like election campaigns and deliberative forums, and who learns the most. The obvious measure, the gain in the share of knowledge items answered correctly, is a poor indicator of true learning. We show, in a linear model, a nonlinear one and a general one, that when those who know more learn more, as the literature on prior knowledge suggests, post-process knowledge has nonnegative covariance with true learning under the general model, overall and conditional on true initial knowledge. Observed gain has no such unconditional guarantee and can be negatively correlated with learning. Conditioning on noisy observed initial knowledge does not inherit the guarantee. In 1,212 simulated processes whose observable features match 21 Deliberative Polls, observed gain’s median correlation with true learning is .01 and post-process knowledge’s is .49. The choice matters: in the same polls, the above-median education difference in observed gain is uncertain (.017 [-.005, .041]), while their post-process knowledge is higher conditional on initial knowledge (.066 [.049, .085]).

1 Introduction

Many communication processes are broadly informative, exposing their participants to some goodly quantity and wide variety of information on at least one sizable topic, like policy or electoral choices. The category includes both natural processes like election campaigns and ad hoc but naturalistic ones like Deliberative Polls. The information exposure need not be the point. Often, it is not. In election campaigns, it is incidental, a by-product of persuasion or mobilization. In Deliberative Polls and other deliberative fora, it is intrinsic but instrumental—not an end in itself but an ingredient in more fully considered, better founded opinions. But in whatever fashion, for whatever reason, the exposure must occur, and at least some participants must learn something from it, or the process would be only degenerately “informative.”

Almost always, we are interested in this learning—sometimes in the learning of this or that particular bit, affecting this or that particular attitude (Gilens 2001), but more often in the amount of topic-wide learning, affecting a wider range of attitudes, as well as vote intentions, levels of engagement, and the like (Bartels 1996; Delli Carpini and Keeter 1996; Althaus 2003). In the aggregate, this is also normatively important, speaking to the educative value of election campaigns and the deliberative quality of deliberative fora.

But how best to measure it? The obvious choice is observed knowledge gain: the increase in the proportion of knowledge items answered correctly. For aggregate description, *faute de mieux*, this may serve. But individual-level analyses—of learning as a function of causally prior variables or of causally subsequent ones as a function of learning—are another story. There, observed knowledge gain turns out to be a surprisingly weak indicator of true knowledge gain. This paper shows

why in two models, a linear one and a nonlinear one, then generalizes the key results to any measurement function, gives conditions for post-process knowledge to track learning and compares the proxies under specified learning models, measures how much better in simulations calibrated to real Deliberative Polls, and shows that the choice changes substantive answers.

2 Observed and True Knowledge

The knowledge we are used to thinking about is “observed.” It is, by definition, all we ever see. Conventionally, the i th person’s observed knowledge at time t —call it x_{it} —is the proportion of the questionnaire’s knowledge items the person answers correctly. We are much less used to thinking about “true knowledge”—what observed knowledge is trying to observe. Its depths can be left to philosophy and neuroscience. Nearer the surface, we may define the i th person’s true knowledge at time t —call it X_{it} —as the proportion of all knowable, relevant items he or she knows at time t . Observed knowledge gain is $\Delta x_i = x_{i2} - x_{i1}$ and true knowledge gain $\Delta X_i = X_{i2} - X_{i1}$.

The set of knowable, relevant items includes everything that is relevant and knowable, whether any mortal knows it or not; on topics like policy issues, it must number at least in the hundreds of thousands. So X_{it} never gets much above 0, and since very few people follow politics very closely, its distribution must be severely right-skewed. By contrast, x_{it} typically ranges all the way from 0 to 1, because the questionnaire’s items are a very small, very easy subset of what is knowable. Question writers do not know enough to ask the items known only by policy experts, let alone those no one knows. Of those they know, they favor relatively easy ones; otherwise, hardly anyone would score much above 0, except by lucky guessing.

These facts suggest the following properties of knowledge and learning:

- L1. The participants learn: $E(X_{i2}|X_{i1}) > X_{i1}$ for $X_{i1} > 0$.
- L2. The more knowledgeable learn more: $E(\Delta X_i|X_{i1})$ increases with X_{i1} . More new information sticks when there is more previously cognized information for it to stick to (Recht and Leslie 1988; Eckhardt et al. 1991; Cooke et al. 1993; Hambrick 2003).
- L3. Those who know nothing learn nothing.
- M1. Those who know nothing show no knowledge, and those who know the most show as much as the measure permits.
- M2. The more knowledgeable appear to know more, but the less so, the more they know: $E(x_{it}|X_{it})$ increases with X_{it} but flattens as it approaches its ceiling of 1.
- M3. People appear to know more than they do: $E(x_{it}|X_{it}) > X_{it}$.
- M4. The more knowledgeable appear to learn more, then, beyond a point, to learn less. Those who know nothing actually learn nothing; those who know the most, despite actually learning the most, cannot show it, having started at or near the ceiling.

Three phenomena are jointly responsible for the results that follow.

Item-sampling bias. A questionnaire’s knowledge items are hardly a random draw from the universe of relevant, knowable items. They are much, much, much easier than average. Question writers cannot even think of the items known only by policy experts, much less those nobody knows. Among those they can think of, they tend to choose items that anyone paying a modicum of attention to the news should know. Otherwise, hardly anyone would score much above 0, except by lucky guessing. Thus many of the relatively knowledgeable may ace a questionnaire’s knowledge items, letting x_{it} range far higher than X_{it} , usually all the way up to 1.

Ceiling effects. The participants who already know a lot, having a relatively high X_{i1} and thus an absolutely high x_{i1} , cannot show much observed knowledge gain. For true knowledge, there is always room for improvement; for observed knowledge, not so.

The rich get richer. Knowing more about something makes it easier to learn still more (Dooling and Lachman 1971; Bransford and Johnson 1972; Recht and Leslie 1988; Eckhardt et al. 1991; Cooke et al. 1993; Hambrick 2003). Learning is easier when the material falls effortlessly into context. Hence the participants who cannot show much observed knowledge gain are precisely those whose true knowledge gain should be greatest.

These considerations yield two second-blush intuitions. First, ceiling effects and the rich getting richer should depress the correlation between Δx_i and ΔX_i . Second, x_{i2} may be a better indicator. Those who emerge observably knowing a lot have generally gained a lot—either observably, if they initially knew only a little, or unobservably, if they initially knew a lot.

3 A Linear Model

Let

$$x_{i1} = a_1 + b_1 X_{i1} + u_{i1}, \quad (1)$$

$$x_{i2} = a_2 + b_2 X_{i2} + u_{i2}, \quad (2)$$

$$X_{i2} = \alpha + \beta X_{i1} + v_i, \quad (3)$$

implying

$$\Delta x_i = (a_2 - a_1) + b_2 \alpha + (b_2 \beta - b_1) X_{i1} + (u_{i2} - u_{i1}) + b_2 v_i, \quad (4)$$

$$\Delta X_i = \alpha + (\beta - 1) X_{i1} + v_i. \quad (5)$$

The measurement errors u_{i1} and u_{i2} and the disturbance v_i have zero means, are mutually uncorrelated, and are uncorrelated with X_{i1} . The measurement errors are also uncorrelated with X_{i2} , and $\sigma_{X_1}^2$, $\sigma_{u_1}^2$, $\sigma_{u_2}^2$ and σ_v^2 are all strictly positive: there is some measurement error at each wave, true initial knowledge does not perfectly predict true later knowledge, and true initial knowledge varies.

A second layer of assumptions expresses the three phenomena. *Item-sampling bias.* Those who know nothing, and do not guess, score 0, so $a_t = 0$ and $b_t = E(x_{it}|X_{it} = X_{Mt})/X_{Mt}$, where X_{Mt} is the highest true knowledge in the population. Since the questionnaire lets $E(x_{it}|X_{it} = X_{Mt})$ approach 1 while X_{Mt} is near 0,

$$b_1, b_2 > 1. \quad (6)$$

The rich get richer. A given difference in true initial knowledge results in a bigger difference in true later knowledge:

$$\beta > 1. \quad (7)$$

Ceiling effects. The more one knows, the less one can appear to gain: X_{i1} 's coefficient in (4) is negative,

$$b_1 > b_2 \beta. \quad (8)$$

The correlations of observed gain and of post-process knowledge with true gain are $\rho_{\Delta x \Delta X} = \sigma_{\Delta x \Delta X} / (\sigma_{\Delta x} \sigma_{\Delta X})$ and $\rho_{x_2 \Delta X} = \sigma_{x_2 \Delta X} / (\sigma_{x_2} \sigma_{\Delta X})$, where single-subscripted, squared σ 's are vari-

ances and double-subscripted ones covariances. From (1)–(5),

$$\sigma_{\Delta x \Delta X} = (\beta - 1)(b_2\beta - b_1)\sigma_{X_1}^2 + b_2\sigma_v^2, \quad (9)$$

$$\sigma_{x_2 \Delta X} = (\beta - 1)b_2\beta\sigma_{X_1}^2 + b_2\sigma_v^2, \quad (10)$$

$$\sigma_{\Delta x}^2 = (b_2\beta - b_1)^2\sigma_{X_1}^2 + \sigma_{u_1}^2 + \sigma_{u_2}^2 + b_2^2\sigma_v^2, \quad (11)$$

$$\sigma_{x_2}^2 = b_2^2\beta^2\sigma_{X_1}^2 + b_2^2\sigma_v^2 + \sigma_{u_2}^2. \quad (12)$$

Three results follow (proofs in Appendix A).

Result 1. $\rho_{\Delta x \Delta X}$ can be negative. Specifically,

- Condition 1a: $\rho_{\Delta x \Delta X} < 0$ if and only if $(\beta - 1)(b_1 - b_2\beta)\sigma_{X_1}^2 - b_2\sigma_v^2 > 0$.

Result 2. $\rho_{x_2 \Delta X}$ is always positive.

Result 3. $\rho_{\Delta x \Delta X}$ can be less than $\rho_{x_2 \Delta X}$. Each of the following is sufficient:

- Condition 3a: $(b_1 - 2b_2\beta)b_1\sigma_{X_1}^2 + \sigma_{u_1}^2 > 0$;
- Condition 3b: $b_1 \geq 2b_2\beta$;
- Condition 3c: Condition 1a;
- Condition 3d: $\gamma_{21} < .5$, where $\gamma_{21} = \sigma_{x_1 x_2} / \sigma_{x_1}^2$ is the slope of the population regression of x_{i2} on x_{i1} .

Three points need noting. First, $\rho_{\Delta x \Delta X}$ need not be negative to be depressed: as the parameters approach a configuration under which it is negative, observed gain is enfeebled, positively but weakly correlated with true gain. Second, none of Conditions 3a–3d is necessary. Third, Condition 3b implies 3a, but Conditions 3a and 3c do not imply one another; neither is uniformly easier to satisfy. Condition 3d is directly checkable: in every one of the 21 Deliberative Polls we examine below, the observed correlation of x_{i1} and x_{i2} is between .11 and .72, and the regression slope of x_{i2} on x_{i1} is below .5 in 10 of them.

Guessing. Let people who do not know an item answer it correctly with probability κr_i , where r_i is the person's propensity to guess, the same at both waves and independent of the other variables, and κ the chance of a lucky guess. Observed knowledge becomes $y_{it} = x_{it}(1 - \kappa r_i) + \kappa r_i$. Then $\sigma_{\Delta y \Delta X} = (1 - \kappa E(r_i))\sigma_{\Delta x \Delta X}$ and $\sigma_{y_2 \Delta X} = (1 - \kappa E(r_i))\sigma_{x_2 \Delta X}$: guessing scales both covariances by the same nonnegative factor. Results 1 and 2 retain their strict signs when the factor is positive; guessing correctly with probability one makes the observed score constant.

4 A Nonlinear Model

The knowledge proportions x_{i1} and x_{i2} are bounded by 0 and 1 and often approach the ceiling; their difference lies between -1 and 1. Their consequently decreasing slopes on true knowledge (M2's ceiling) are a major part of the story, which the linear model captures only through $b_1 > b_2\beta$. A quadratic measurement equation makes the ceiling explicit. Let

$$X_{i2} = \beta X_{i1}, \quad (13)$$

$$x_{it} = bX_{it} + cX_{it}^2 + u_{it}, \quad t = 1, 2, \quad (14)$$

with $\beta > 1$, $b > 0$ and $c < 0$, so that

$$\Delta X_i = (\beta - 1)X_{i1}, \quad (15)$$

$$\Delta x_i = (\beta - 1)bX_{i1} + (\beta^2 - 1)cX_{i1}^2 + u_{i2} - u_{i1}. \quad (16)$$

M2 requires observed knowledge to rise with true knowledge over the range of true knowledge at both waves: $b + 2cX_{it} > 0$, or $X_{M2} = \beta X_{M1} \leq -b/2c$. $E(\Delta x_i | X_{i1})$, by contrast, peaks at $X_{i1} = -b/2(\beta + 1)c$, below X_{M1} only if the support extends beyond that point. M4 therefore requires $X_{M1} > -b/[2(\beta + 1)c]$ as well as $\beta > 1$. Writing V and C for variances and covariances,

$$C(\Delta x_i, \Delta X_i) = (\beta - 1)^2 [bV(X_{i1}) + (\beta + 1)cC(X_{i1}^2, X_{i1})], \quad (17)$$

$$C(x_{i2}, \Delta X_i) = (\beta - 1)\beta [bV(X_{i1}) + c\beta C(X_{i1}^2, X_{i1})], \quad (18)$$

$$V(\Delta X_i) = (\beta - 1)^2 V(X_{i1}), \quad (19)$$

$$V(\Delta x_i) = (\beta - 1)^2 b^2 V(X_{i1}) + (\beta^2 - 1)^2 c^2 V(X_{i1}^2) + 2bc(\beta - 1)(\beta^2 - 1)C(X_{i1}^2, X_{i1}) + V(u_{i1}) + V(u_{i2}), \quad (20)$$

$$V(x_{i2}) = b^2 \beta^2 V(X_{i1}) + c^2 \beta^4 V(X_{i1}^2) + 2bc\beta^3 C(X_{i1}^2, X_{i1}) + V(u_{i2}). \quad (21)$$

Since $X_{i1} \geq 0$, $C(X_{i1}^2, X_{i1}) > 0$, and with $c < 0$ the second term in each bracket is negative. Let $\gamma = C(X_{i1}^2, X_{i1})/V(X_{i1})$, the slope of the population regression of X_{i1}^2 on X_{i1} . Then:

- R1. $\rho_{\Delta x \Delta X} > 0$ if and only if $\gamma < -b/(\beta + 1)c$.
- R2. $\rho_{x_2 \Delta X} > 0$ if and only if $\gamma < -b/\beta c$.
- R3. Since $-b/(\beta + 1)c < -b/\beta c$, $\rho_{\Delta x \Delta X} > 0$ implies $\rho_{x_2 \Delta X} > 0$, but not the reverse.
- R4. Under M2, R2's condition always holds: $\rho_{x_2 \Delta X} > 0$. The bracket in (18) is the covariance of X_{i1} with $bX_{i1} + c\beta X_{i1}^2$, a function that rises over the whole range of X_{i1} when $\beta X_{M1} \leq -b/2c$; the covariance of a variable with an increasing function of itself is positive. The bracket in (17) is the covariance of X_{i1} with $bX_{i1} + (\beta + 1)cX_{i1}^2$, which rises and then falls, and can be negative.

5 How General Are These Results?

Both models are special cases of a general one. Let each questionnaire item j have a difficulty d_j on the scale of true knowledge, and let a person know the item with probability $F(X_{it}, d_j)$, nondecreasing in X_{it} . In the simplest case F is a threshold—the person knows the item if and only if $X_{it} > d_j$ —and observed knowledge is $x_{it} = G(X_{it})$, where G is the distribution function of the questionnaire's difficulties. An easy questionnaire concentrates its difficulties low, so G rises quickly and flattens at 1. Evenly spread difficulties give a piecewise linear G ; a smooth, concentrated distribution gives something close to the quadratic; a logistic F gives an item response model. Let people who do not know an item answer it correctly by lucky guess, and let true learning be $\Delta X_i = h(X_{i1}, R_i)$, where R_i is a scalar learning shock independent of X_{i1} , and h is nondecreasing in R_i .

Proposition 1. *If h is nondecreasing in X_{i1} (the rich get richer), $\text{Cov}(x_{i2}, \Delta X_i) \geq 0$.*

Proposition 2. *For any h , $\text{Cov}(x_{i2}, \Delta X_i | X_{i1}) \geq 0$.*

Both follow from the association of independent random variables: any two nondecreasing functions of independent variables have nonnegative covariance (Harris 1960; Esary et al. 1967). Nothing in either argument depends on the shape of G , the distribution of true knowledge, or the amount of guessing (Appendix A). Observed gain has no such guarantee: $\Delta x_i = G(X_{i1} + \Delta X_i) - G(X_{i1})$ rises and then falls in X_{i1} when G has a ceiling, while true gain rises throughout. With deterministic measurement, conditioning on X_{i1} fixes x_{i1} , so Δx_i and x_{i2} differ only by a constant. Conditioning on the realized observed baseline also makes gain and post-process scores equivalent, but does

not condition on true initial knowledge. The two distinctions matter when baseline knowledge is measured with error (Lord 1967; Allison 1990).

6 How Large Are the Differences?

The propositions say nothing about magnitudes. To gauge them, we simulate processes and ask how well each proxy tracks true gain. Each simulated process has 300 participants, like a Deliberative Poll. The simulated true-knowledge variable is a right-skewed gamma trait, rather than a literal proportion of all knowable items, and item difficulties are drawn from a range of its quantiles, so that questionnaires run from very easy to moderately hard. People know items either by threshold or with logistic probability, and guess correctly with probability 0, .1 or .25. We vary the number of items (5 to 15), the median learning-rate multiplier and its log-scale dispersion, and three forms of learning: proportional to initial knowledge (L2), unrelated to it, and larger for those who know less.

Within the specified parameter ranges, we retain processes whose observable summaries resemble the polls. For each of 21 Deliberative Polls in the corrected, source-based dp-data export of the polls in Luskin et al. (2022), we compute five observable features: the mean and standard deviation of initial observed knowledge, the mean observed gain, the correlation of initial and post-process knowledge, and the correlation of observed gain with initial knowledge. All five empirical features use participants with knowledge scores at both waves. A simulated process passes the filter if each feature falls within the corresponding range across polls; it need not resemble any single poll jointly. Inputs are frozen at dp-data commit 681d08ad72dfb50fa24eccfb837fd48f3db9a700, with SHA-256 checksums in the repository manifest. Of 4,000 processes with proportional learning, 30% are plausible; with learning unrelated to initial knowledge, 6%; with catch-up learning, 8 of 4,000 processes. Proportional learning passes more often in this grid, but these rates are conditional on the chosen model family and parameter distribution. They do not establish which learning process generated the polls. The gain–initial-knowledge correlation is negative in every poll (-.66 to -.35). Ceilings can produce this pattern, but so can baseline measurement error, because the baseline score is subtracted in observed gain.




Figure 1: Correlation of each proxy with true knowledge gain across simulated processes whose observable features fall within the marginal ranges across Deliberative Polls. Points are medians and lines span the 5th to 95th percentiles across processes; they describe variation across plausible worlds, not sampling uncertainty. Post-process knowledge given initial is the partial correlation holding observed initial knowledge constant.

Figure 1 shows the results. When the knowledgeable learn more, observed gain’s median correlation with true gain is .01, negative in 46% of plausible processes; post-process knowledge’s is .49, never negative, and higher than observed gain’s in 100% of them. Holding observed initial knowledge constant, post-process knowledge’s partial correlation is .34, higher than observed gain’s in 100% of processes. When learning is unrelated to initial knowledge, the three proxies are similar (medians .16, .19 and .21); observed gain performs best in the 8 retained catch-up processes. This small selected subset does not characterize catch-up learning generally.

Could observed gain be rehabilitated? Asking more items helps only slightly. Among plausible processes with proportional learning, observed gain’s median correlation is .00 with 5 items and .03 with 15, while post-process knowledge’s is .45 and .55. Asking harder items reduces the ceiling but can only be pushed so far: question writers do not know many of the hardest items, and legitimately shy away from asking the hardest they do know, lest nearly everyone score zero. Scoring by item response theory (Jackman 2000; Levendusky and Jackman 2003) or weighting by difficulty (Sood 2011) changes little, since such scores correlate above .96 with conventional ones (Luskin and Bullock 2006). Branching, as in computerized adaptive testing, has more promise, but item pools are capped by question writers’ expertise and respondents’ patience (Montgomery and Cutler 2013), and no scoring can recover items no one knew to ask.

7 Does the Choice Matter?

Consider a question the literature on deliberation asks often: who learns? In the same 21 polls, we regress each learning proxy on whether education is strictly above the poll’s participant median, poll by poll, and pool the estimates in a Bayesian random-effects meta-analysis (Röver 2020). With observed gain, the above-median education difference is uncertain (.017 [-.005, .041]), with a negative estimate in 6 of the polls. With post-process knowledge given initial knowledge, the above-median group scores higher (.066 [.049, .085]), with a negative estimate in 1 polls (Figure 2). Post-process knowledge alone, without initial knowledge, gives .116 [.098, .134], mixing what people learned with what they already knew. As the model implies, observed gain with initial knowledge held constant gives exactly the same estimate as post-process knowledge with initial knowledge held constant.

The education indicator comes from dp-data’s reviewed poll-specific ordered measures. Its fixed reference population is the unique historical-polardata participant sample with observed education, before exclusions for outcomes or group IDs; ties at the median remain in the lower group. We use the exported indicator without recomputing cutoffs. Education, paired knowledge and discussion-group membership must be observed for a participant to enter the regressions. The comparisons are associations with observed scores, not causal effects of education or direct measurements of true learning. Poll estimates use HC1 standard errors clustered by discussion group and normal 95% intervals. The meta-analysis uses a half-normal prior with scale .1 on between-poll heterogeneity and an improper uniform prior on the pooled mean; its intervals are 95% posterior credible intervals.




Figure 2: Difference in learning proxies for participants above versus at or below their poll's education median, by poll, with 95% intervals. Left: observed gain. Right: post-process knowledge holding initial knowledge constant. Standard errors are clustered by small group; the pooled estimates are from Bayesian random-effects meta-analyses.

Learning is often also used to explain attitude change: in Deliberative Polls, the participants who learn the most have been expected to account for much of the net change in opinion. We estimate these associations on all 129 attitude indices in the corrected data, regressing the change in each index on a learning proxy and the participant's distance from the other complete-case members of the small group at baseline. Records with missing scores or group IDs and groups with fewer than two complete cases are excluded. The net direction and proxy standard deviation are computed within that same analysis sample. The share of indices on which the coefficient is positive is 50% for observed gain, 59% for post-process knowledge and 53% for post-process knowledge given initial knowledge; the shares significantly positive are 5%, 16% and 12%, respectively (Appendix Table 2). These summaries show proxy-specific associations; they do not establish that learning has no relationship with attitude change. Initial knowledge also has positive associations. The tests are unadjusted for multiple comparisons, indices within a poll are dependent, and the direction is estimated from the same sample. The significance shares therefore are descriptive rather than a

calibrated test against chance.

8 Conclusion

Together, the oversampling of easy items, easily neared or hit ceilings, and the tendency of the most knowledgeable to learn the most make observed knowledge gain a weak measure of true knowledge gain. In these simulations, post-process knowledge, with baseline adjustment where appropriate, is often a more informative proxy. This is a measurement comparison under the stated assumptions, not identification of a causal effect of learning. When learning is nondecreasing in true initial knowledge and an independent scalar learning shock, the general model guarantees nonnegative covariance between post-process knowledge and learning. The corresponding conditional result holds given true initial knowledge; it does not guarantee a positive partial correlation given noisy observed initial knowledge. The simulations illustrate the distinction, and the poll comparisons show how conclusions about observed learning proxies depend on adjustment and measurement.

References

- Allison, Paul D. 1990. “Change Scores as Dependent Variables in Regression Analysis.” *Sociological Methodology* 20: 93–114. <https://doi.org/10.2307/271083>.
- Althaus, Scott L. 2003. *Collective Preferences in Democratic Politics: Opinion Surveys and the Will of the People*. Cambridge University Press. <https://doi.org/10.1017/cbo9780511610042>.
- Bartels, Larry M. 1996. “Uninformed Votes: Information Effects in Presidential Elections.” *American Journal of Political Science* 40 (1): 194–230. <https://doi.org/10.2307/2111700>.
- Bransford, John D., and Marcia K. Johnson. 1972. “Contextual Prerequisites for Understanding: Some Investigations of Comprehension and Recall.” *Journal of Verbal Learning and Verbal Behavior* 11 (6): 717–26. [https://doi.org/10.1016/s0022-5371\(72\)80006-9](https://doi.org/10.1016/s0022-5371(72)80006-9).
- Cooke, Nancy J., Robert S. Atlas, David M. Lane, and Robert C. Berger. 1993. “Role of High-Level Knowledge in Memory for Chess Positions.” *The American Journal of Psychology* 106 (3): 321–51. <https://doi.org/10.2307/1423181>.
- Delli Carpini, Michael X., and Scott Keeter. 1996. *What Americans Know about Politics and Why It Matters*. Yale University Press.
- Dooling, D. James, and Roy Lachman. 1971. “Effects of Comprehension on Retention of Prose.” *Journal of Experimental Psychology* 88 (2): 216–22. <https://doi.org/10.1037/h0030904>.
- Eckhardt, Beverly B., Mary R. Wood, and Robin Smith Jacobvitz. 1991. “Verbal Ability and Prior Knowledge: Contributions to Adults’ Comprehension of Television.” *Communication Research* 18 (5): 636–49. <https://doi.org/10.1177/009365091018005004>.
- Esary, J. D., F. Proschan, and D. W. Walkup. 1967. “Association of Random Variables, with Applications.” *The Annals of Mathematical Statistics* 38 (5): 1466–74. <https://doi.org/10.1214/aoms/1177698701>.

- Gilens, Martin. 2001. “Political Ignorance and Collective Policy Preferences.” *American Political Science Review* 95 (2): 379–96. <https://doi.org/10.1017/s0003055401002222>.
- Hambrick, David Z. 2003. “Why Are Some People More Knowledgeable Than Others? A Longitudinal Study of Knowledge Acquisition.” *Memory & Cognition* 31 (6): 902–17. <https://doi.org/10.3758/bf03196444>.
- Harris, T. E. 1960. “A Lower Bound for the Critical Probability in a Certain Percolation Process.” *Mathematical Proceedings of the Cambridge Philosophical Society* 56 (1): 13–20. <https://doi.org/10.1017/s0305004100034241>.
- Jackman, Simon. 2000. “Estimation and Inference Are Missing Data Problems: Unifying Social Science Statistics via Bayesian Simulation.” *Political Analysis* 8 (4): 307–32. <https://doi.org/10.1093/oxfordjournals.pan.a029818>.
- Levendusky, Matthew S., and Simon D. Jackman. 2003. “Reconsidering the Measurement of Political Knowledge.” December. https://web.stanford.edu/class/polisci353/2004spring/reading/levendusky_final.pdf.
- Lord, Frederic M. 1967. “A Paradox in the Interpretation of Group Comparisons.” *Psychological Bulletin* 68 (5): 304–5. <https://doi.org/10.1037/h0025105>.
- Luskin, Robert C., and John G. Bullock. 2006. “Measuring Political Knowledge.” Unpublished manuscript.
- Luskin, Robert C., Gaurav Sood, James S. Fishkin, and Kyu S. Hahn. 2022. “Deliberative Distortions? Homogenization, Polarization, and Domination in Small Group Discussions.” *British Journal of Political Science* 52 (3): 1205–25. <https://doi.org/10.1017/s0007123421000168>.
- Montgomery, Jacob M., and Josh Cutler. 2013. “Computerized Adaptive Testing for Public Opinion Surveys.” *Political Analysis* 21 (2): 172–92. <https://doi.org/10.1093/pan/mps060>.
- Recht, Donna R., and Lauren Leslie. 1988. “Effect of Prior Knowledge on Good and Poor Readers’ Memory of Text.” *Journal of Educational Psychology* 80 (1): 16–20. <https://doi.org/10.1037/0022-0663.80.1.16>.
- Röver, Christian. 2020. “Bayesian Random-Effects Meta-Analysis Using the Bayesmeta r Package.” *Journal of Statistical Software* 93 (6). <https://doi.org/10.18637/jss.v093.i06>.
- Sood, Gaurav. 2011. “A More Discriminating Way to Score: Guessing and Hardness Adjusted Scoring.” Unpublished manuscript.

A Proofs

Result 1 and Condition 1a. Since $\rho_{\Delta x \Delta X}$ ’s sign is that of $\sigma_{\Delta x \Delta X}$, $\rho_{\Delta x \Delta X} < 0$ exactly when $(\beta - 1)(b_2\beta - b_1)\sigma_{X_1}^2 + b_2\sigma_v^2 < 0$. The second addend is always positive; by (7) and (8), the first is always negative. So $\rho_{\Delta x \Delta X} < 0$ if and only if $-(\beta - 1)(b_2\beta - b_1)\sigma_{X_1}^2$ exceeds $b_2\sigma_v^2$, which is

Condition 1a.

Result 2. $\sigma_{x_2\Delta X} = (\beta - 1)b_2\beta\sigma_{X_1}^2 + b_2\sigma_v^2$, and both addends are positive by (6) and (7).

Result 3 and Conditions 3a–3d. $\rho_{\Delta x\Delta X} < \rho_{x_2\Delta X}$ if and only if $\sigma_{x_2\Delta X}/\sigma_{x_2} > \sigma_{\Delta x\Delta X}/\sigma_{\Delta x}$. A jointly sufficient condition is that $\sigma_{x_2\Delta X} > \sigma_{\Delta x\Delta X}$ and $\sigma_{x_2}^2 < \sigma_{\Delta x}^2$. By (9) and (10), the first is equivalent, after subtracting $b_2\sigma_v^2$ and dividing by $(\beta - 1)\sigma_{X_1}^2$, to $b_2\beta > b_2\beta - b_1$, that is, $b_1 > 0$, which always holds. So $\sigma_{x_2}^2 < \sigma_{\Delta x}^2$ alone is sufficient. By (11) and (12), it holds if and only if $b_2^2\beta^2\sigma_{X_1}^2 < (b_2\beta - b_1)^2\sigma_{X_1}^2 + \sigma_{u_1}^2$, that is, $(b_1 - 2b_2\beta)b_1\sigma_{X_1}^2 + \sigma_{u_1}^2 > 0$: Condition 3a. Condition 3b makes the first term nonnegative. Condition 3c follows from Results 1 and 2: if $\rho_{\Delta x\Delta X} < 0 < \rho_{x_2\Delta X}$, the inequality holds. For 3d, since $\Delta x_i = x_{i2} - x_{i1}$, $\sigma_{x_2}^2 < \sigma_{\Delta x}^2$ if and only if $0 < \sigma_{x_1}^2 - 2\sigma_{x_1x_2}$, that is, $\gamma_{21} < .5$.

Guessing. With $k_i = \kappa r_i$, $\Delta y_i = (1 - k_i)\Delta x_i$ and $y_{i2} = (1 - k_i)x_{i2} + k_i$. Since k_i is independent of Δx_i , x_{i2} and ΔX_i , $C((1 - k_i)\Delta x_i, \Delta X_i) = E(1 - k_i)C(\Delta x_i, \Delta X_i)$, and likewise for y_{i2} .

Propositions 1 and 2. Condition on the item difficulties. Let U_i collect the independent randomness in knowing items and in guessing at time 2. Then $x_{i2} = \phi(X_{i1} + h(X_{i1}, R_i), U_i)$ with ϕ nondecreasing in its first argument. Write $\bar{\phi}(X_{i1}, R_i) = E[x_{i2} | X_{i1}, R_i]$, nondecreasing in both arguments when h is. Since ΔX_i does not depend on U_i , $\text{Cov}(x_{i2}, \Delta X_i) = \text{Cov}(\bar{\phi}(X_{i1}, R_i), h(X_{i1}, R_i))$. X_{i1} and R_i are independent, hence associated (Esary et al. 1967), and two nondecreasing functions of associated variables have nonnegative covariance. Given X_{i1} , $\bar{\phi}$ and h are nondecreasing functions of R_i alone, whose covariance is nonnegative (Harris 1960). Proposition 2 concerns true initial knowledge; holding observed initial knowledge constant, a coarse and noisy measure of it, the partial correlation was negative in 1 of the 1,212 plausible simulated processes with proportional learning.

B The polls and further results

Table 1: Observable features of the Deliberative Polls. The first five judge which simulated processes are plausible; γ_{21} is the slope of the regression of x_2 on x_1 (Condition 3d).

Poll	Mean x_1	Mean gain	SD x_1	$r(x_1, x_2)$	$r(\Delta x, x_1)$	γ_{21}
San Mateo, CA	.15	.17	.14	.11	-.46	.18
Europolis	.20	.17	.16	.21	-.46	.29
Central Power & Light	.29	.16	.19	.27	-.62	.26
Zeguo Township	.31	.11	.31	.30	-.50	.35
By the People: Health and Education	.35	.06	.21	.35	-.55	.36
Tomorrow’s Europe (EU)	.37	.18	.19	.47	-.44	.52
New Haven, CT	.37	.23	.18	.39	-.35	.53
National Issues Convention 2	.37	.12	.17	.56	-.44	.58
By the People: National	.39	.03	.18	.61	-.45	.60
West Texas Utilities	.40	.29	.24	.39	-.56	.38
Bulgarian National	.43	.09	.25	.29	-.58	.30
Southwestern Electric Power	.44	.20	.23	.22	-.57	.25
National Issues Convention	.48	.10	.20	.62	-.59	.52
Australia Republic Referendum	.48	.28	.23	.59	-.66	.46
By the People 2004 US Presidential Primaries	.51	.12	.28	.72	-.42	.69
UK EU	.52	.20	.29	.45	-.52	.45
UK Crime	.52	.09	.23	.61	-.36	.67
UK General Election	.54	.10	.21	.53	-.52	.51
UK Monarchy	.65	.15	.19	.56	-.42	.60
UK Health	.66	.08	.23	.56	-.42	.59
By the People 2004 US General Election	.75	.07	.24	.71	-.51	.63

Table 2: Learning and attitude change. Each attitude index’s change is regressed on the proxy and the distance from the small group’s mean; the proxy’s coefficient is signed so that positive means moving with the sample’s net change. Significance at the 5% level with standard errors clustered by small group.

Proxy	Indices	Positive	Significantly positive	Significantly negative
Observed gain	129	.50	.05	.05
Post-process knowledge	129	.59	.16	.06
Post-process knowledge given initial	129	.53	.12	.04
Initial knowledge	129	.59	.10	.09